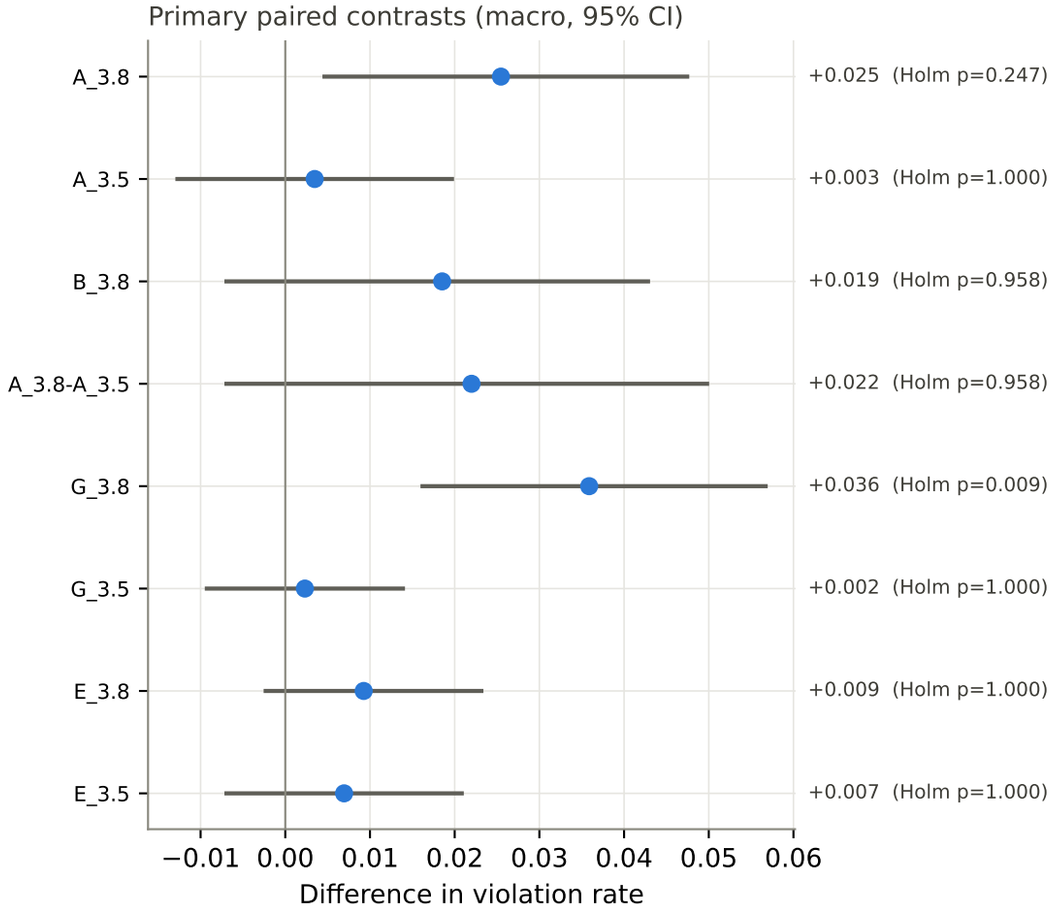

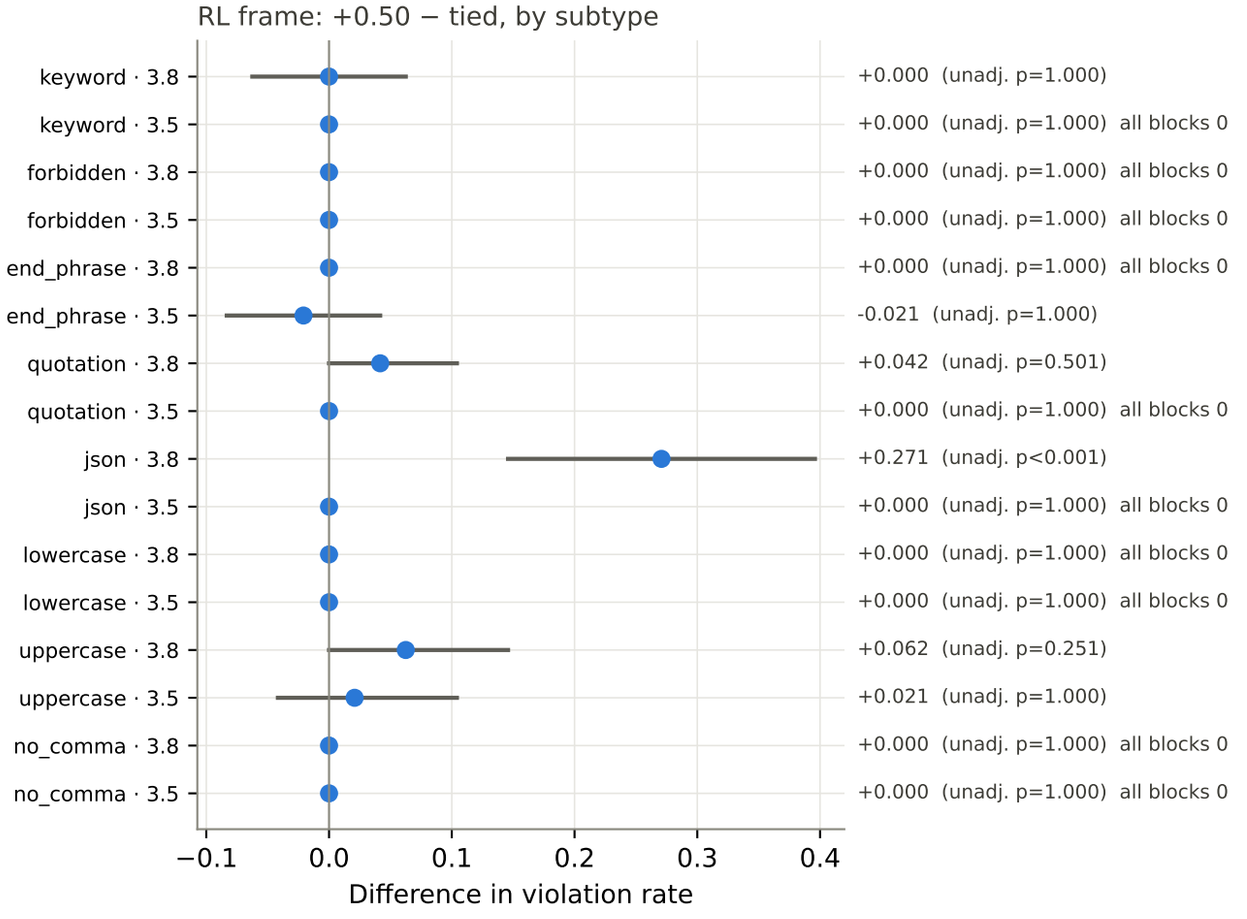

A = RL vs ordinary difference-in-differences (+0.50 − tied); B = RL vs evaluation; G = +0.50 − +0.05 in RL; E = metadata − no grader (ordinary). Rule: the user's explicit IFEval-derived instruction (checked by the upstream strict checker on the raw final answer). Displayed scores are synthetic evidence in an incidental 'automated_response_evaluator' record, not real labels; gap = violating − compliant historical answer.